

## Analyzing volatile memory on a Google Kubernetes Engine node

Marcus Hallberg  
Security engineer - Detection & Response  
Spotify



## Disclaimer

- This talk presents personal research and applications of open source tools.
- Spotify is not liable for this content or published code.
- Questions, remarks or other actions should be aimed towards the presenter.



## Agenda

- Main takeaways
- Container workloads at Spotify
- Google Kubernetes Engine
- What are memory snapshots?
- How can we create a memory snapshot on GKE?
- Demo and analysis

## Main takeaways

- Accessing and analysing volatile memory on a GKE node.
- Complete setup using open source tools.
- All code used - <https://github.com/Monrava/cyberthreat2023>



# Container workloads at Spotify



**CYBER THREAT  
2023**

**5 GCP Regions**

**4096 nodes/cluster**

**+4000 services**

**+3000 namespaces**



## The problem:

- We get an alert that a user accessed a container in GKE.
- What did user do once they accessed the container?
- There are no commercial tools available.



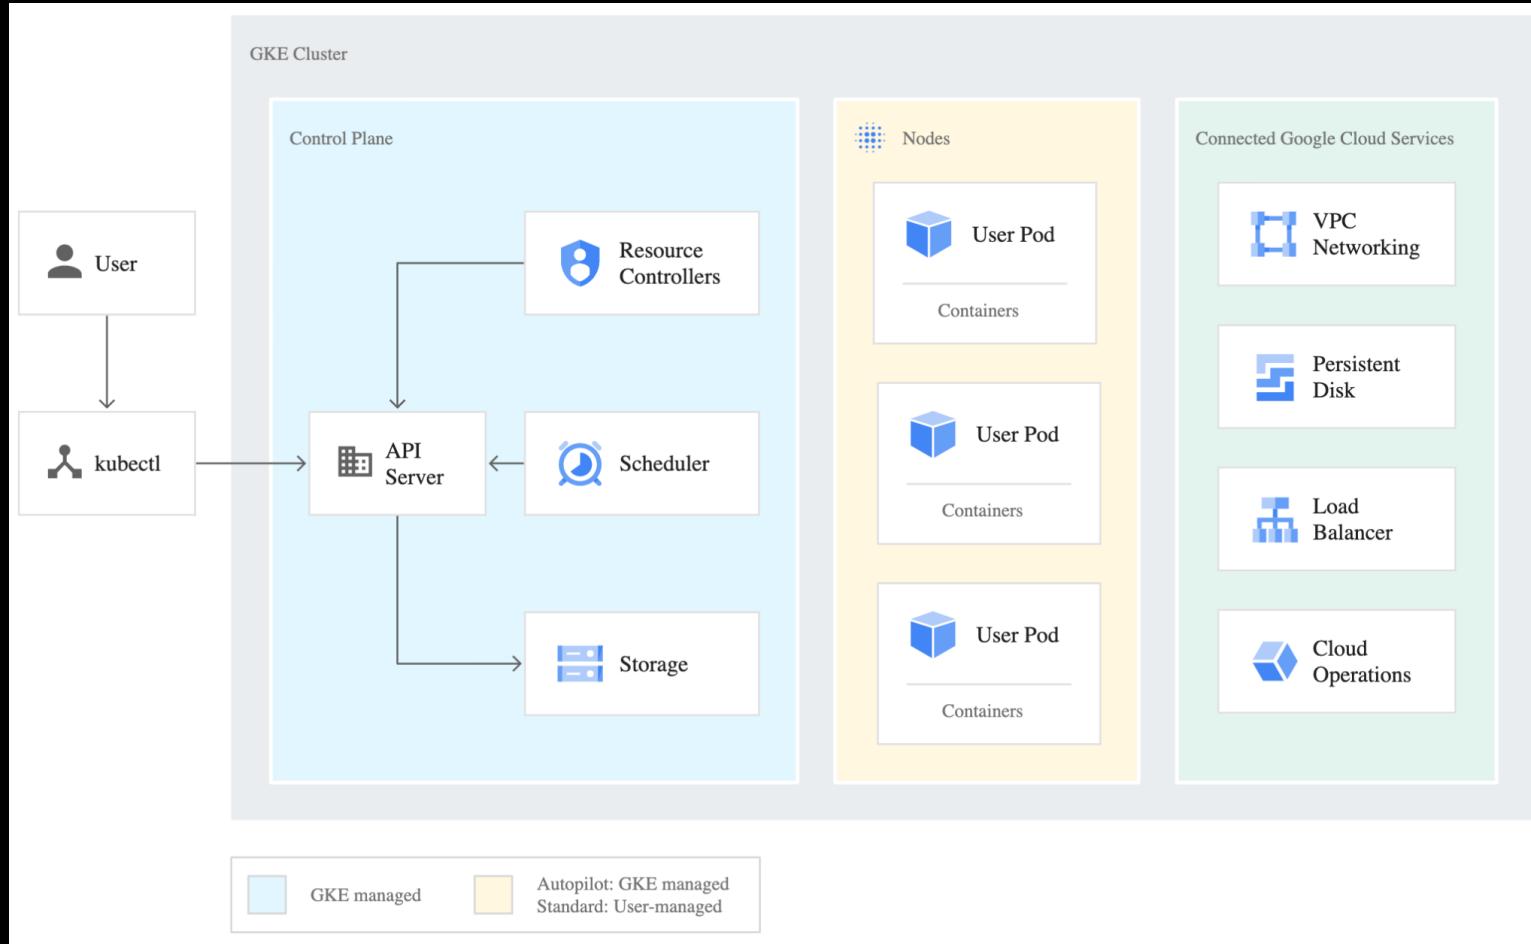


```
"principalEmail": "user@spotify.com"  
  "granted": true,  
  "permission": "io.k8s.core.v1.pods.exec.create",  
  "resource": "core/v1/namespaces/default/pods/pod-node-affinity-zeroday-attacker-pod/exec"  
  "callerIp": "195.178.178.197",  
  "callerSuppliedUserAgent": "kubectl/v1.25.4 (darwin/amd64; client-go/1.25.4; pkg/0.4.0)",  
  "resourceName": "core/v1/namespaces/default/pods/pod-node-affinity-zeroday-attacker-pod/exec",  
  "cluster_name": "zeroday-gke-cluster",  
  "project_id": "PROJECT_NAME",  
  "location": "europe-west1-b",  
  "timestamp": "2023-04-11T12:23:43.423331Z",  
  "logName": "projects/PROJECT_NAME/logs/cloudaudit.googleapis.com",  
  "id": "d42f3772-56da-4658-b996-3eaf024b6b59",  
  "receiveTimestamp": "2023-04-11T12:23:57.235797324Z"
```

# Google Kubernetes Engine



# CYBERTHREAT 2023



# What are memory snapshots?

## Memory snapshots

- What are they?
- Are they still used?
- Why would we take memory snapshots on GKE nodes?



## The problem:

- We get an alert that a user accessed a container in GKE.
- What did user do once they accessed the container?
- There are no commercial tools available.



**How can we create a  
memory snapshot on GKE?**

## How do we create a GKE memory snapshot?

- Step 1 - Kernel memory dump:
  - /proc/kcore
  - AVML
- Step 2 - Symbol file:
  - vmlinux
  - dwarf2json
- Step 3 - Analysis (Volatility3)

## The How: Step 1 - Kernel memory dump

- /proc/kcore

*This file represents the physical memory of the system and is stored in the ELF core file format.*

*The total length of the file is the size of physical memory (RAM) plus 4 KiB.*

- AVML (Or other tool able to take dumps of this file)



## The How: Step 1 - Kernel memory dump

- How do we access /proc/kcore?

```
container {
  name  = "avml-container"
  image = var.container_image_avml
  security_context {
    privileged = "true"
    capabilities {
      add = ["CAP_NET_ADMIN", "CAP_SYS_ADMIN"]
    }
  }
}
```

## The How: Step 2 - Symbol file

- vmlinu - what is it?

[https://storage.googleapis.com/cos-tools/\\$build\\_id/vmlinu](https://storage.googleapis.com/cos-tools/$build_id/vmlinu)

<https://storage.googleapis.com/cos-tools/16919.235.1/vmlinu>

- dwarf2json

Images	
<a href="#">EDIT</a> <a href="#">DELETE</a> <a href="#">CREATE INSTANCE</a> <a href="#">EXPORT</a>	
gke-1249-gke3200-cos-97-16919-235-1-v230120-c-pre	
Location	us (United States)
Architecture	x86/64
Labels	None
Creation time	Feb 2, 2023, 11:44:14 AM UTC+01:00
Encryption type	Google-managed

## The How: Step 3 - Analysis

- Volatility3

**volatilityfoundation/  
volatility3**

Volatility 3.0 development



46

Contributors

83

Used by

2k

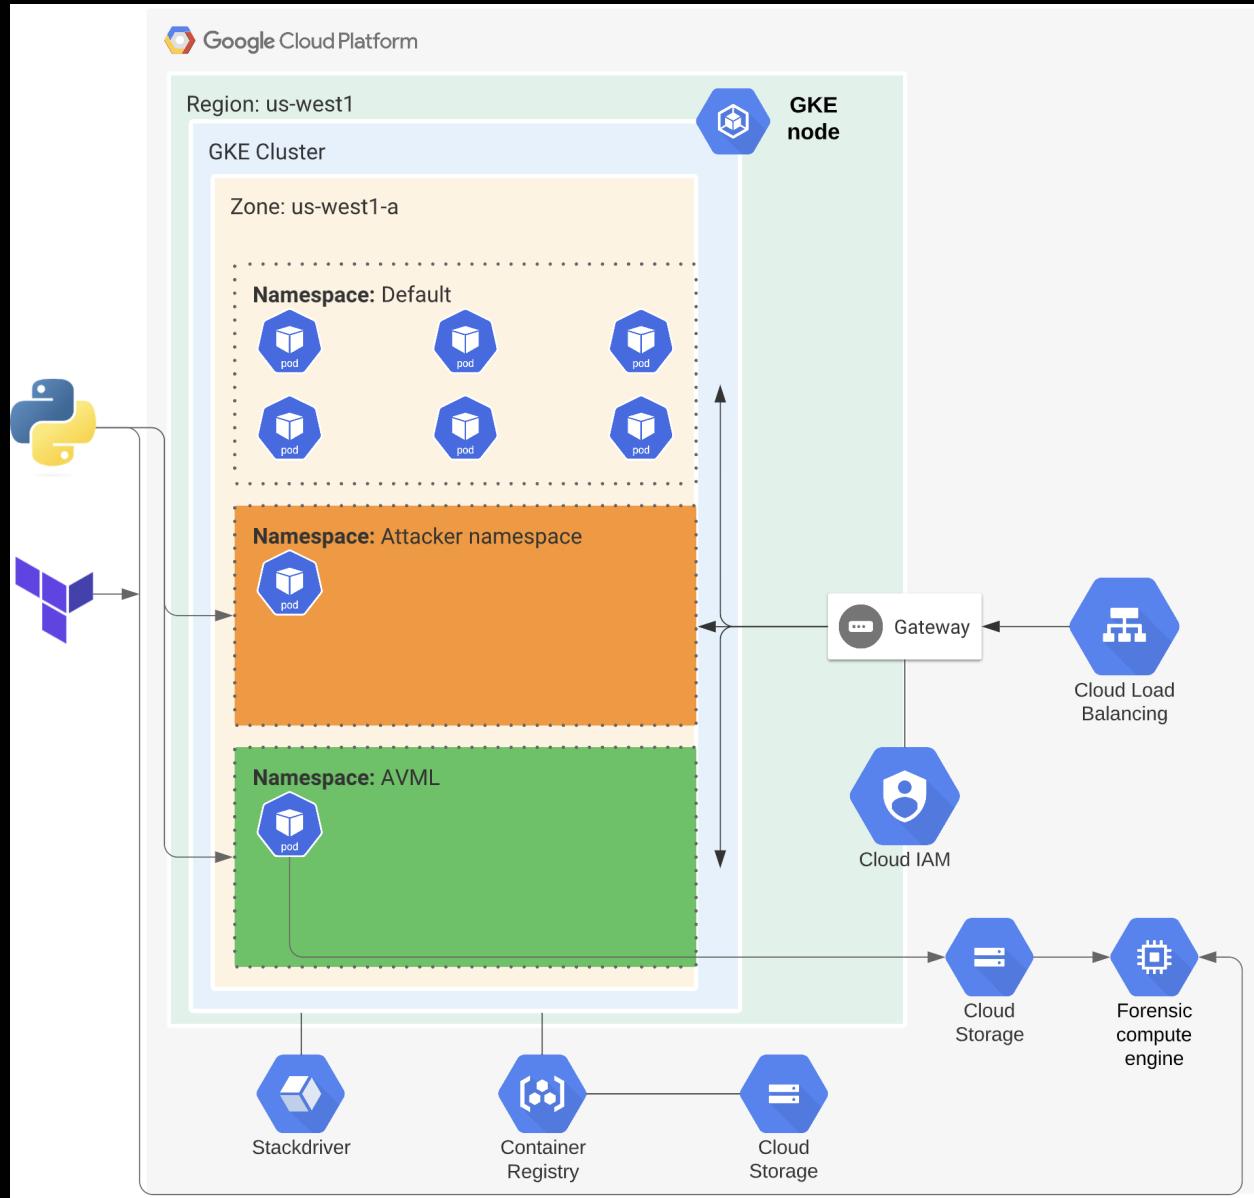
Stars

321

Forks



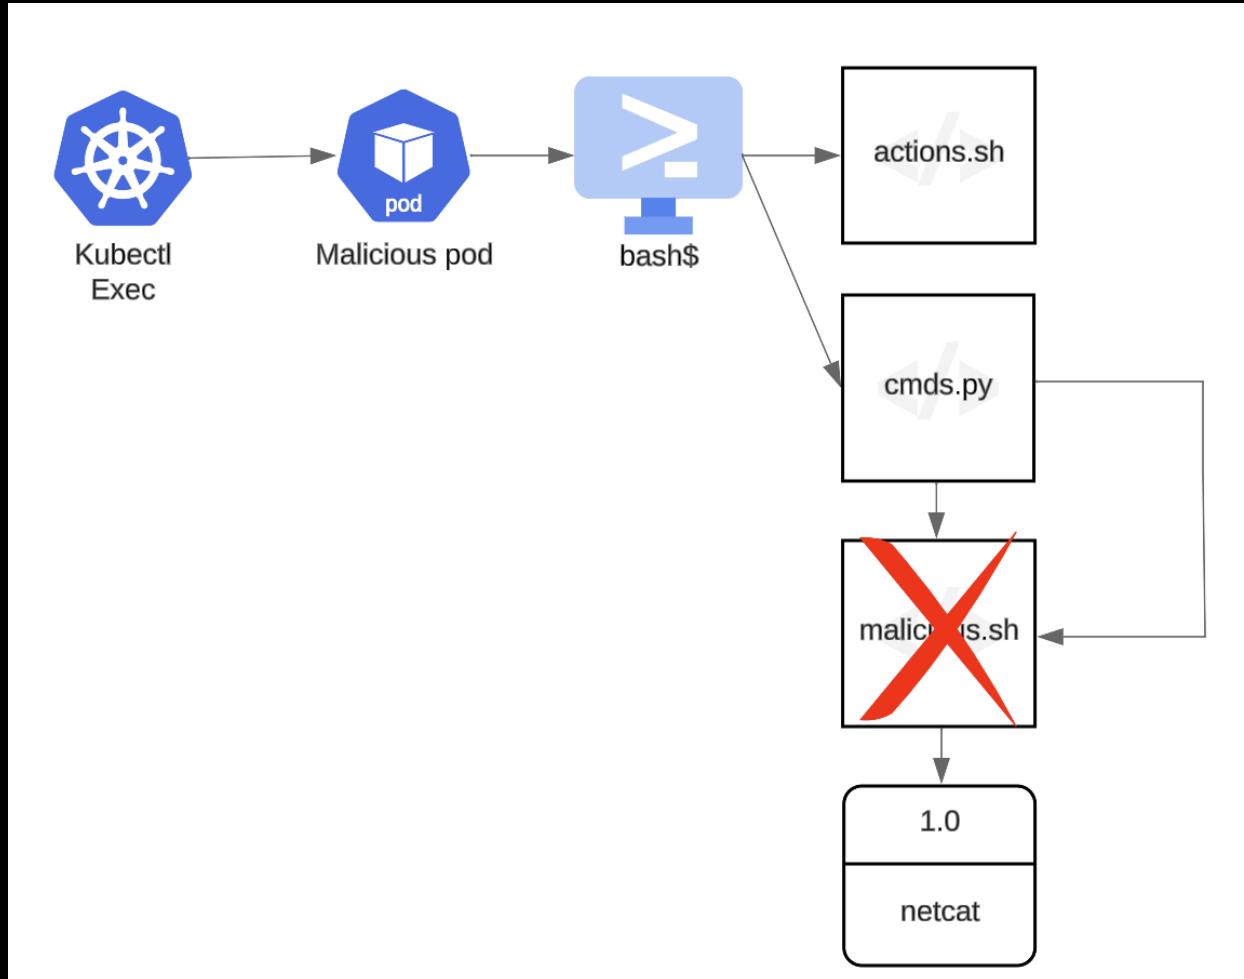
# Testing setup



# Test and analysis

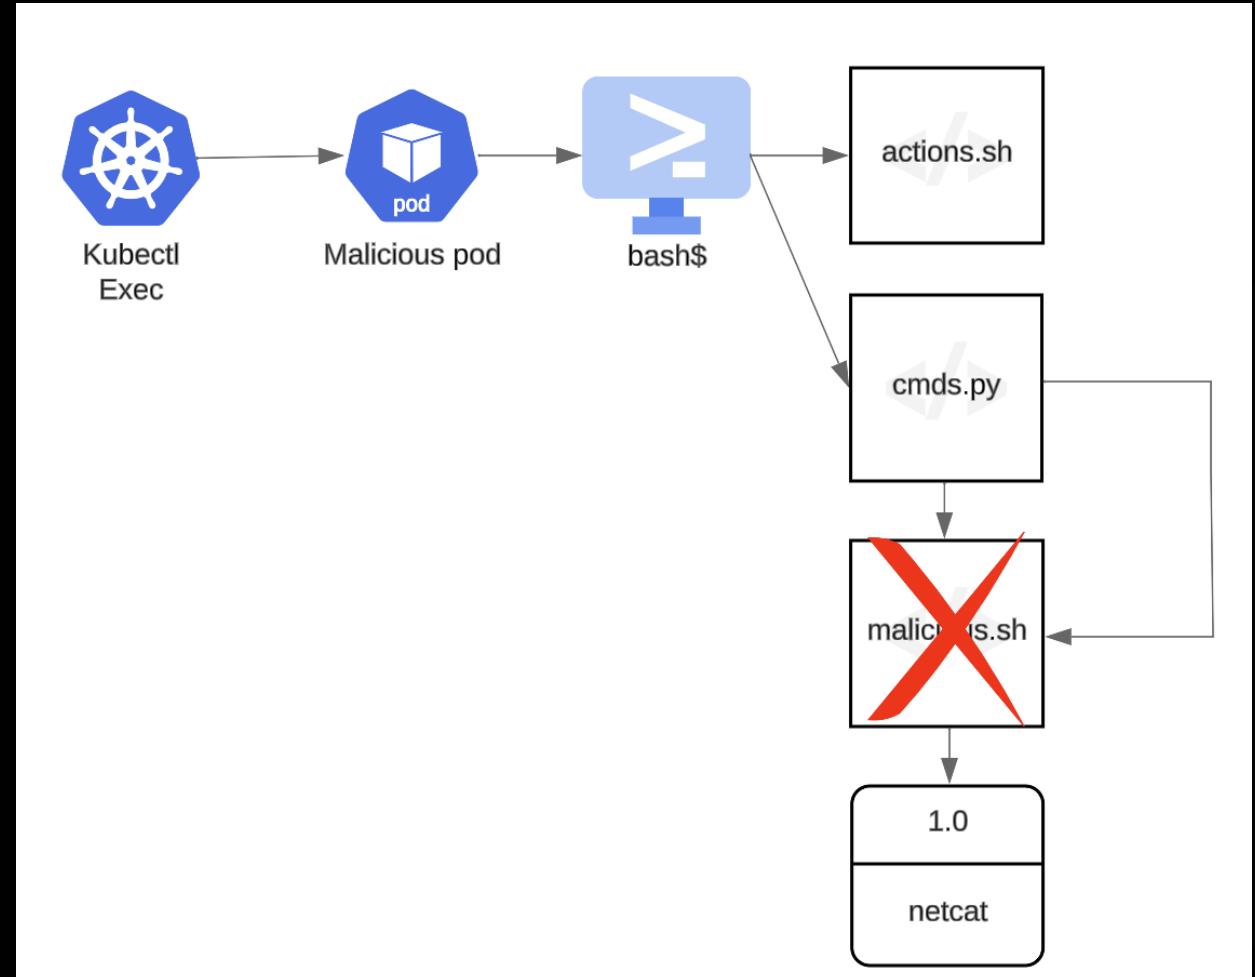
## What we tested

- Various commands using bash
  - Running processes
  - Checking external IP address
  - Running scripts
- Started processes
  - Netcat
  - Python3
  - Bash shell script





## What can we see from the GKE memory dump?





PID	Process	CommandTime	Command
82141	bash	2023-11-16 12:33:43.000000	ls -la /proc/
82141	bash	2023-11-16 12:33:43.000000	unshare -Urm
82141	bash	2023-11-16 12:33:43.000000	echo \$SHELL
82141	bash	2023-11-16 12:33:43.000000	cat /proc/1/cgroup
82141	bash	2023-11-16 12:33:43.000000	HISTFILE=~/.bash_history
82141	bash	2023-11-16 12:33:43.000000	
82141	bash	2023-11-16 12:33:43.000000	#1700138023
82141	bash	2023-11-16 12:33:43.000000	#!/bin/bash -i
82141	bash	2023-11-16 12:33:43.000000	ps aux
82141	bash	2023-11-16 12:33:43.000000	nohup watch -n1 curl -s ifconfig.co &
82141	bash	2023-11-16 12:33:43.000000	grep Cap /proc/1/status
82141	bash	2023-11-16 12:33:43.000000	set -o history
82141	bash	2023-11-16 12:33:43.000000	whoami
82141	bash	2023-11-16 12:33:43.000000	ls -la
82141	bash	2023-11-16 12:33:43.000000	history
82141	bash	2023-11-16 12:33:43.000000	./actions.sh
82141	bash	2023-11-16 12:33:43.000000	ls /root
82141	bash	2023-11-16 12:33:43.000000	pwd
82141	bash	2023-11-16 12:33:43.000000	ls -la
82141	bash	2023-11-16 12:33:43.000000	df -h
82141	bash	2023-11-16 12:33:43.000000	exit

Bash history plugin (Commands are captured in memory)

0x9e2601848fc0	58495	58495	58054	watch
0x9e2603f94ec0	58531	58531	57701	bash
0x9e2528de1f80	58687	58687	58054	python3
0x9e260ace1f80	58754	58754	58687	sh
0x9e260ace2f40	58756	58756	58754	nc
0x9e260ace5e80	58763	58763	58687	bash
0x9e25018d5e80	58839	58839	57892	avml

PSList plugin (Both container activities)



# CYBERTHREAT 2023

```
*** 0x9e2601848fc0  58495  58495  58054  watch
*** 0x9e2528de1f80  58687  58687  58054  python3
**** 0x9e260ace1f80  58754  58754  58687  sh
***** 0x9e260ace2f40   58756   58756   58754  nc
```

PSTree plugin (Attacker container)

PSaux plugin (Both container activities)

```
58054  watch  watch -n1 curl -s ifconfig.co
57701  bash   /bin/bash
58054  python3 python3 cmd.py
58687  sh    /bin/sh -c ./malicious_script.sh
58754  nc    /bin/nc -l 10514
58687  bash   [bash]
```

# Lessons learned

## What have we learned?

- AVML + Dwarf2json + Volatility3
  - NodeSelector
  - Acquire vmlinu for GKE/COS using build\_id
  - Be aware of updates to kernel or compiler
- Avoid overutilization
  - Cluster rebalancing - evicted
  - ContainerStatusUnknown with exit code 137



- Get in touch
  - Tw/X: @monrava89
  - LinkedIn: in/hallbergmarcus
  - @: [mhallberg@spotify.com](mailto:mhallberg@spotify.com)
- Code
  - <https://github.com/Monrava/cyberthreat2023>



**THANK YOU!**